

LC 101

Unit 3 - Regular Expressions

April 10, 2017

Regular Expressions

- A *regular expression* is a pattern that describes a set of strings
 - (Often abbreviated as *re* or *regex*)
- Most often used to test if a candidate string matches a pattern

Basics

- A regular expression is like a mathematical formula
 - But built from characters and three basic operators: concatenation, alternation, and Kleene closure (the * operator)
 - Can also use parentheses to group parts

Concatenation and Alternation

- Individual characters are valid regular expressions
 - **a** is a regular expression that matches exactly one string: “**a**”
- Multi-character sequences formed by *concatenation* are valid regular expressions
 - **ab** is a regex that matches exactly one string: “**ab**”
- Alternation is the “or” operator. It can be used to select between two (or more) subexpressions
 - **a|b** is a regex that matches exactly two strings: “**a**” and “**b**”
 - **a|b|c** matches three strings: “**a**”, “**b**”, and “**c**”

Concatenation and Alternation

- The operands of the alternation or concatenation operators can be any valid subexpressions
- Concatenation takes precedence over alternation
- Parentheses can be used to override precedence
 - `ab | c` matches the strings “`ab`” and “`c`”
 - `a(b | c)` matches the strings “`ab`” and “`ac`”
 - `a(b | cd | e) f` matches the strings “`abf`”, “`acdf`”, and “`aef`”

Kleene Closure

- The *Kleene closure* (or *star* operator) is a unary operator that means zero or more repetitions of the preceding subexpression
 - Creates an infinite set of matching strings
 - a^* matches the strings "", "a", "aa", "aaa", "aaaa", "aaaaa", ...
 - Has higher precedence than concatenation
 - ab^* matches "a", "ab", "abb", "abbb", "abbbb", ...
 - $(ab)^*$ matches "", "ab", "abab", "ababab", "abababab", ...
 - $(a|b)^*$ matches "", "a", "b", "aa", "ab", "ba", "bb", "aaa", "aab", "aba", "abb", "baa", ...

Other Operations

- Although those three basic operations (concatenation, alternation, and Kleene closure) are all that are needed to create any expression, a few other helpful operations are usually defined
 - The plus operator means one or more of the preceding subexpression
 - **a+** matches the strings “a”, “aa”, “aaa”, “aaaa”, ...
 - The question mark operator means zero or one of the preceding subexpression
 - **a?** matches “” and “a”
 - **ab?c** matches “ac” and “abc”
 - **{n}** means exactly *n* iterations of the preceding subexpression
 - **ab{3}c** matches the string “abbbc”
 - **{m,n}** means at least *m* but no more than *n* of the preceding subexpression
 - **ab{1,2}c** matches the strings “abc” and “abbc”

Limitations

- While powerful, regular expressions are limited in the types of sets they can describe
 - There are seemingly simple cases that regular expressions cannot handle
 - For example, there is no regular expression that can be used to determine if a string has properly balanced nested parentheses (without additional restrictions)
 - This means that we cannot use a regular expression to determine if a string is a valid regular expression!
 - Or mathematical expression

From Theory to Practice

- Many programming languages have regular expressions as either part of the language itself or as part of their standard libraries
- Though there is a POSIX standard, most languages and libraries seem to have their own slightly different variant
 - The theory is the same but the exact syntax varies

Any Character

- A dot matches any single valid character
 - Includes special characters but sometimes does not include the newline character
 - . matches “a”, “b”, ..., “1”, “2”, ..., “%”, “@”, ...

Character Sets

- A character set matches any character contained in the brackets
 - `[abc]` matches “a”, “b”, and “c”
- Character sets can use a dash to indicate a range of characters
 - `[a-e]` is the same as `[abcde]`
- A character set can be negated by starting it with the hat character
 - `[^abc]` means any character *except* a, b, or c

Special Characters

- `\d` matches a digit character. Equivalent to `[0-9]`
- `\D` matches a non-digit character. Equivalent to `[^0-9]`
- `\s` matches a whitespace character
 - Space, tab, form feed, line feed
- `\S` matches a non-whitespace character
- `\t` matches a tab
- `\n` matches a form feed character
- `\r` matches a carriage return
- `\w` matches an alphanumeric character or underscore. Equivalent to `[A-Za-z0-9_]`
- `\W` matches a non-word character. Equivalent to `[^A-Za-z0-9_]`

Escaping Special Characters

- The backslash `\` is used to escape special characters
 - `*` matches the literal “`*`” and does not mean the Kleene closure
 - `\.` matches a literal “`.`” and does not mean any character
- Must also be used to escape a literal backslash
 - `\\s` matches “`\s`”

Creating Regular Expressions

- Regular Expressions can be created using either a regular expression literal or the RegExp constructor
 - The literals are compiled when the script is loaded, so use them for better performance for constant expressions
 - The RegExp constructor allows specification of an expression at runtime

```
var re1 = /ab*/;
```

```
var re2 = new RegExp('ab*');
```

Finding Matches

- `re.test(str)` searches for a match in a string and returns `true` or `false`
- `str.search(re)` searches for a match in a string and returns the index of the first match found
 - And returns -1 if no match is found
- Note that these search for a match in a string. They do not check that the entire string matches.
- Also note that `test` is a method of `RegExp` while `search` is a method of `String`

```
var re = /ab*/;  
var str = 'cabbc';  
console.log(re.test(str));    // outputs true  
console.log(str.search(re));  // outputs 1
```

Start and End of Input

- Often we want to see if an entire string matches an expression and not just if it contains a match
- Special characters can be used in regular expressions to match the start or end of input
 - `^` matches the start of input and `$` matches the end of input
 - So using `^` at the beginning of an expression and `$` at the end will then require the entire string to match

```
var re = /^ab*$/;
var str1 = 'abb';
var str2 = 'cabbc';
console.log(re.test(str1)); // outputs true
console.log(re.test(str2)); // outputs false
```


Spaces

- Note that the literal space character matters in regular expressions!

```
var re = /hello goodbye/;
var str1 = 'hello goodbye';
var str2 = 'hellogoodbye';
console.log(re.test(str1)); // outputs true
console.log(re.test(str2)); // outputs false
```

Capturing Groups

- When parentheses are used in regular expressions, you can access the substring that matched the subexpression in the parentheses
 - `re.exec(str)` returns a result array on successful match or `null` if not found
 - `str.match(re)` returns a result array on successful match or `null` if not found

```
var re = /^(ab*)(c*d)$/;
```

```
var str = 'abbccd';
```

```
console.log(re.exec(str));
```

```
// outputs ["abbccd", "abb", "ccd", index: 0, input: "abbccd"]
```

```
console.log(str.match(re));
```

```
// outputs ["abbccd", "abb", "ccd", index: 0, input: "abbccd"]
```

Non-Capturing Parentheses

- If you need to use parentheses but don't care about the subgroup then you can use `(?:x)`

```
var re = /^(?:ab*)(c*d)$/;
```

```
var str = 'abbccd';
```

```
console.log(re.exec(str));
```

```
// outputs ["abbccd", "ccd", index: 0, input: "abbccd"]
```

Greedy vs Lazy

- The *, +, and ? operators are greedy
 - The match as much as they can to get a valid match
- You can make the operators lazy by putting another ? after them
 - Then they will match as little as they can to get a valid match

```
var str = 'abcbc';
```

```
var re1 = /a.*c/;    // greedy
console.log(re1.exec(str));
// outputs ["abcbc", index: 0, input: "abcbc"]
```

```
var re2 = /a.*?c/;   // lazy
console.log(re2.exec(str));
// outputs ["abc", index: 0, input: "abcbc"]
```

Replace

- The `string.replace` method can be used to replace parts of a string that match a regular expression
 - By default it will just replace the first match. To replace all matches, the *global modifier* must be used on the regular expression
 - The modifiers come after the closing slash in a regexp literal or as the second parameter of the `RegExp` constructor

```
var re = /fish/g;
var str = 'red fish blue fish';
console.log(str.replace(re, 'bird')); // outputs "red bird blue bird"
```

Multiple Matches

- The global modifier also changes how match and exec work
 - match will return a simple array of the strings that match instead of the first match
 - exec can be called multiple times to return subsequent matches

```
var str = 'red fish blue fish';  
console.log(str.match(/fish/)); // outputs ["fish", index: 4, input: "red fish blue fish"]  
console.log(str.match(/fish/g)); // outputs ["fish", "fish"]
```

```
var re = /fish/g;  
console.log(re.exec(str)); // outputs ["fish", index: 4, input: "red fish blue fish"]  
console.log(re.exec(str)); // outputs ["fish", index: 14, input: "red fish blue fish"]
```

Real-Life Example

- In JavaScript, there is no existing function to determine if an entire string is an int
 - `parseInt` will return a number if the first part of a string is digits
 - The `Number` constructor will work for both floats and ints

```
var str1 = '42a';
```

```
console.log(parseInt(str1, 10)); // outputs 42
```

```
var str2 = '4.2';
```

```
console.log(new Number(str2)); // outputs Number {[[PrimitiveValue]]: 4.2}
```

Real-Life Example

- We can use a regular expression to determine if a string contains only digits
 - What would happen if the expression below was just “/**d**+”?

```
var intRegExp = /^d+$/;
```

```
console.log(intRegExp.test('42a')); // outputs false
```

```
console.log(intRegExp.test('4.2')); // outputs false
```

```
console.log(intRegExp.test('42')); // outputs true
```